

[< Back to AI Programming with Python Nanodegree](#)

# Image Classifier Application

REVIEW

CODE REVIEW 8

HISTORY

▼ train.py 4

```
1 #!/usr/bin/env python3
2 # -*- coding: utf-8 -*-
3 # *AIPND/ImageClassifierApp/train.py
4 #
5 # trains a new network on a dataset of images
6 # specs :
7 # The training loss, validation loss, and validation accuracy are printed out :
8 # allows users to choose from at least two different architectures available for
9 # allows users to set hyperparameters for learning rate, number of hidden units
10 # allows users to choose training the model on a GPU
11 #
12 # Example calls:
13 # Ex 1, use data_dir 'flowers': python train.py flowers
14 # Ex 2, use save_dir 'chksav' to save checkpoint: python train.py --save_dir cl
15 # Ex 3, use densenet161 and hidden_units '1000, 500': python train.py --arch de
16 # Ex 4, set epochs to 10: python train.py -e 10
17 # Ex 5, set learning rate to 0.002: python train.py -lr 0.002
18 # Ex 6, train in GPU mode (subject to device capability): python train.py --gpu
19
20 import argparse
21 import torch
22 import numpy as np
23 from torch import nn
24 from torch import optim
25 import torch.nn.functional as F
26 from torchvision import datasets, transforms, models
27 from PIL import Image
```

```

28
29 from datetime import datetime
30 import os
31 import glob
32 import sys
33
34 from workspace_utils import active_session
35
36 device = torch.device('cuda:0' if torch.cuda.is_available() else 'cpu')
37 model_names = ['densenet121', 'densenet161', 'vgg16']
38 datadir = 'flowers'
39 savedir = 'chksav'
40
41 # main
42 def main():
43     # get input arguments and print
44     args = get_input_args()
45     print('\n*** command line arguments ***')
46     print('architecture:', args.arch, '\ndata dir:', args.data_dir, '\nchkpt d:
47         '\nlearning rate:', args.learning_rate, '\nhidden layer:', args.hidd
48         '\nepochs:', args.epochs, '\nGPU mode:', args.gpu, '\n')

```



AWESOME

Well done clearly logging the validation loss and accuracy at each step!

```

49
50 if len(glob.glob(args.data_dir)) == 0:
51     print('*** data dir: ', args.data_dir, ', not found ... exiting\n')
52     sys.exit(1)
53
54 if args.learning_rate <= 0:
55     print('*** learning rate cannot be negative or 0 ... exiting\n')
56     sys.exit(1)
57
58 # if hidden units supplied, check values are numeric
59 if args.hidden_units:
60     try:
61         list(map(int, args.hidden_units.split(',')))
62     except ValueError:
63         print("hidden units contain non numeric value(s) :", args.hidden_u
64         sys.exit(1)
65
66 if args.epochs < 1:
67     print('*** epochs cannot be less than 1 ... exiting\n')
68     sys.exit(1)
69
70 # transform and load training, validation and testing sets
71 dataloaders = transform_load(args)
72
73 # load pre-trained model and replace with custom classifier
74 model = models.__dict__[args.arch](pretrained=True)
75 model = build_classifier(model, args, dataloaders)
76 print('\n*** model architecture:', args.arch, '\n*** Classifier:\n', model.c
77
78 # set training criterion and optimizer
79 criterion = nn.NLLLoss()
80 optimizer = optim.Adam(model.classifier.parameters(), args.learning_rate)

```

```

80
81 # start model training and testing
82 if device.type == 'cuda':
83     if args.gpu:
84         print('*** GPU is available, using GPU ...\n')
85     else:
86         print('*** training classifier in GPU mode ...\n')
87 else:
88     if args.gpu:
89         print('*** GPU is unavailable, using CPU ...\n')
90     else:
91         print('*** training classifier in CPU mode ...\n')
92
93 with active_session():
94     model = train(model, dataloaders, optimizer, criterion, args.epochs, 4)
95     model = test(model, dataloaders, criterion)
96
97 # save to checkpoint
98 model = model.cpu() # back to CPU mode post training
99 model.class_to_idx = dataloaders['train'].dataset.class_to_idx
100
101 # if checkpoint dir not exists, create it
102 if not os.path.isdir(args.save_dir):
103     os.makedirs(args.save_dir)
104
105 checkpoint = {
106     'classifier': model.classifier,
107     'state_dict': model.state_dict(),
108     'class_to_idx': model.class_to_idx,
109     'optimizer': optimizer.state_dict(),
110     'arch': args.arch,
111     'lr': args.learning_rate,
112     'epochs': args.epochs}
113
114 chkpt = datetime.now().strftime('%Y%m%d_%H%M%S') + '_' + args.arch + '.pth
115 checkpoint = os.path.join(args.save_dir, chkpt)
116
117 torch.save(checkpoint, checkpoint)
118 print('\n*** checkpoint: ', chkpt, ', saved to: ', os.path.dirname(checkpoint))
119
120
121 def get_input_args():
122     # create parser
123     parser = argparse.ArgumentParser()
124
125     parser.add_argument('data_dir', type=str, nargs='?', default=datadir,
126                        help='path to datasets')
127
128     parser.add_argument('--save_dir', type=str, default=savedir,
129                        help='path to checkpoint directory')
130
131     parser.add_argument('--arch', dest='arch', default='densenet121',
132                        choices=model_names, help='model architecture: ' +
133                        ' | '.join(model_names) + ' (default: densenet121)')
134
135     parser.add_argument('-lr', '--learning_rate', dest='learning_rate', default=
136                        help='learning rate (default: 0.001)')
137
138     parser.add_argument('-hu', '--hidden_units', dest='hidden_units', default=N
139                        help='hidden units, one or multiple values (comma sepa
140

```

```

141         """ enclosed in single quotes. Ex1. one value: '500'
142             Ex2. multiple values: '1000, 500' """
143
144     parser.add_argument('-e', '--epochs', dest='epochs', default=3, type=int,
145                         help='total no. of epochs to run (default: 3)')
146
147     parser.add_argument('--gpu', dest='gpu', default=False, action='store_true'
148                         help='train in gpu mode')

```



AWESOME

Well done declaring the arguments for configuring the command line application!

```

149
150     return parser.parse_args()
151
152 def transform_load(args):
153     train_dir = args.data_dir + '/train'
154     valid_dir = args.data_dir + '/valid'
155     test_dir = args.data_dir + '/test'
156
157     normalize = transforms.Normalize(mean=[0.485, 0.456, 0.406],
158                                     std=[0.229, 0.224, 0.225])
159
160     # define transforms
161     data_transforms = {
162         'train': transforms.Compose([
163             transforms.RandomRotation(30),
164             transforms.RandomResizedCrop(224),
165             transforms.RandomHorizontalFlip(),
166             transforms.ToTensor(),
167             normalize]),
168         'valid': transforms.Compose([
169             transforms.Resize(256),
170             transforms.CenterCrop(224),
171             transforms.ToTensor(),
172             normalize]),
173         'test': transforms.Compose([
174             transforms.Resize(256),
175             transforms.CenterCrop(224),
176             transforms.ToTensor(),
177             normalize])
178     }
179
180     # define datasets
181     image_datasets = {
182         'train': datasets.ImageFolder(train_dir, transform=data_transforms['train']),
183         'valid': datasets.ImageFolder(valid_dir, transform=data_transforms['valid']),
184         'test': datasets.ImageFolder(test_dir, transform=data_transforms['test'])
185     }
186
187     # define dataloaders
188     dataloaders = {
189         'train': torch.utils.data.DataLoader(image_datasets['train'], batch_size=32, shuffle=True),
190         'valid': torch.utils.data.DataLoader(image_datasets['valid'], batch_size=32, shuffle=False),
191         'test': torch.utils.data.DataLoader(image_datasets['test'], batch_size=32, shuffle=False)
192     }

```

```

193     return dataloaders
194
195
196 def build_classifier(model, args, dataloaders):
197     # Freeze parameters so we don't backprop through them
198     for param in model.parameters():
199         param.requires_grad = False
200
201     in_size = {
202         'densenet121': 1024,
203         'densenet161': 2208,
204         'vgg16': 25088,
205     }
206
207     hid_size = {
208         'densenet121': [500],
209         'densenet161': [1000, 500],
210         'vgg16': [4096, 4096, 1000],

```

AWESOME

Great job allowing the user to specify the model architecture between DenseNet121, DenseNet161 and VC

```

211     }
212
213     output_size = len(dataloaders['train'].dataset.classes)
214     relu = nn.ReLU()
215     dropout = nn.Dropout()
216     output = nn.LogSoftmax(dim=1)
217
218     if args.hidden_units:
219         h_list = args.hidden_units.split(',')
220         h_list = list(map(int, h_list)) # convert list from string to int
221     else:
222         h_list = hid_size[args.arch]
223
224     h_layers = [nn.Linear(in_size[args.arch], h_list[0])]
225     h_layers.append(relu)
226     if args.arch[:3] == 'vgg':
227         h_layers.append(dropout)
228
229     if len(h_list) > 1:
230         h_sz = zip(h_list[:-1], h_list[1:])
231         for h1, h2 in h_sz:
232             h_layers.append(nn.Linear(h1, h2))
233             h_layers.append(relu)
234             if args.arch[:3] == 'vgg':
235                 h_layers.append(dropout)
236
237     last = nn.Linear(h_list[-1], output_size)
238     h_layers.append(last)
239     h_layers.append(output)
240
241     print(h_layers)
242     model.classifier = nn.Sequential(*h_layers)
243
244     return model
245
# validate model

```

```

247 def validate(model, dataloaders, criterion):
248     valid_loss = 0
249     accuracy = 0
250
251     for images, labels in iter(dataloaders['valid']):
252
253         images, labels = images.to(device), labels.to(device)
254
255         output = model.forward(images)
256         valid_loss += criterion(output, labels).item()
257         ps = torch.exp(output)
258         equality = (labels.data == ps.max(dim=1)[1])
259         accuracy += equality.type(torch.FloatTensor).mean()
260
261     return valid_loss, accuracy
262
263
264 # train model

```

#### SUGGESTION

Please use `docstrings` to document your major classes and methods. Here's the [reference](#)

```

265 def train(model, dataloaders, optimizer, criterion, epochs=2, print_freq=20, lr=0.01):
266
267     model.to(device)
268     start_time = datetime.now()
269
270     print('epochs:', epochs, ', print_freq:', print_freq, ', lr:', lr, '\n')
271
272     steps = 0
273
274     for e in range(epochs):
275         model.train()
276         running_loss = 0
277         for images, labels in iter(dataloaders['train']):
278             steps += 1
279
280             images, labels = images.to(device), labels.to(device)
281
282             optimizer.zero_grad()
283
284             output = model.forward(images)
285             loss = criterion(output, labels)
286             loss.backward()
287             optimizer.step()
288
289             running_loss += loss.item()
290
291             if steps % print_freq == 0:
292                 model.eval()
293
294                 with torch.no_grad():
295                     valid_loss, accuracy = validate(model, dataloaders, criterion)
296
297                 print('Epoch: {}/{}..'.format(e+1, epochs),
298                       'Training Loss: {:.3f}..'.format(running_loss/print_freq),
299                       'Validation Loss: {:.3f}..'.format(valid_loss/len(dataloaders['valid'])))

```

```

299             'Validation Accuracy: {:.3f}%'.format(accuracy/len(datalo
300             )
301             running_loss = 0
302
303             model.train()
304
305             elapsed = datetime.now() - start_time
306
307             print('\n*** classifier training done ! \nElapsed time[hh:mm:ss.ms]: {}'.fo
308             return model
309
310 # test model
311 def test(model, dataloaders, criterion):
312     print('\n*** validating testset ...\n')
313     model.cpu()
314     model.eval()
315
316     test_loss = 0
317     total = 0
318     match = 0
319
320     start_time = datetime.now()
321
322     with torch.no_grad():
323         for images, labels in iter(dataloaders['test']):
324
325             model, images, labels = model.to(device), images.to(device), labels
326
327             output = model.forward(images)
328             test_loss += criterion(output, labels).item()
329             total += images.shape[0]
330             equality = labels.data == torch.max(output, 1)[1]
331             match += equality.sum().item()
332
333     model.test_accuracy = match/total * 100
334     print('Test Loss: {:.3f}'.format(test_loss/len(dataloaders['test'])),
335           'Test Accuracy: {:.2f}%'.format(model.test_accuracy))
336
337     elapsed = datetime.now() - start_time
338     print('\n*** test validation done ! \nElapsed time[hh:mm:ss.ms] {}:'.format
339     return model
340
341 # Call to main function to run the program
342 if __name__ == "__main__":
343     main()
344
345

```

► predict.py 3

► README.md 1

► workspace\_utils.py

RETURN TO PATH

Rate this review

---

[Student FAQ](#)